

Exercise: Control Flow

Day 4, Part A

```
> library(foreign)
> library(ggplot2)
> library(reshape2)
> library(microbenchmark)
```

1. There is a series of files containing different years of data on income and poverty in Washington state ("data/wa_income_[year]"). Load and combine all of the data in these files into a single data frame. Hint: this requires just one loop combined with some if/else statements.

```
> main_dir <- "C:/Users/ngraetz/Documents/repos/r_training_penn/" # CHANGE TO YOUR LOCAL COPY OF THE DATA
>
> data <- NULL
> for (year in 1997:2015) {
+   if (year < 2004) {
+     sub <- read.csv(paste0(main_dir, "/data/wa_income_", year, ".csv"))
+   } else {
+     sub <- read.dta(paste0(main_dir, "/data/wa_income_", year, ".dta"))
+     sub <- plyr::rename(sub, c("FIPS" = "fips", "median_income" = "income_median"))
+   }
+   data <- rbind(data, sub)
+ }
> summary(data)
```

fips	year	income_median	poverty
Min. :53001	Min. :1997	Min. :27453	Min. : 6.60
1st Qu.:53019	1st Qu.:2001	1st Qu.:36992	1st Qu.:11.50
Median :53039	Median :2006	Median :42369	Median :14.10
Mean :53039	Mean :2006	Mean :43726	Mean :14.26
3rd Qu.:53059	3rd Qu.:2011	3rd Qu.:48693	3rd Qu.:16.40
Max. :53077	Max. :2015	Max. :81816	Max. :32.30

2. Make a line plot of median household income (y-axis) vs year (x-axis) for each county, saving these as separate pages in a PDF. Hint: the unique() function is useful for finding all the unique values of a vector.

```
> pdf(paste0(main_dir, "output/wa_median_income_trends.pdf"), width=10, height=8)
> for (cnty in unique(data$fips)) {
+   gg <- ggplot(data[data$fips == cnty,], aes(x=year, y=income_median)) +
+     geom_line() + labs(title = paste("County:", cnty))
+   print(gg)
+ }
> dev.off()
pdf
2
```

3. Using a loop, calculate the mean poverty rate in each year. Do the same using dcast() and compare your results.

```
> mean_poverty1 <- data.frame(year = 1997:2015, poverty = NA)
> for (yy in 1997:2015) {
+   mean_poverty1[mean_poverty1$year == yy, "poverty"] <-
+     mean(data[data$year == yy, "income_median"])
+ }
```

```

+ }
>
> mean_poverty2 <- dcast(data, year ~ ., value.var = "income_median",
+                       fun.aggregate = mean)
> names(mean_poverty2)[2] <- "poverty"
>
> all.equal(mean_poverty1, mean_poverty2)
[1] TRUE

```

Bonus:

- Using the `microbenchmark()` function in the `microbenchmark` library, determine which of the two approaches in question 3 is faster, and by how much. Which approach do you prefer? Is the difference in timing enough to sway your opinion? Hint: <https://www.r-bloggers.com/5-ways-to-measure-running-time-of-r-code/> has some helpful examples of using `microbenchmark()`.

```

> microbenchmark("loop" = {
+   mean_poverty1 <- data.frame(year = 1997:2015, poverty = NA)
+   for (yy in 1997:2015) {
+     mean_poverty1[mean_poverty1$year == yy, "poverty"] <-
+       mean(data[data$year == yy, "income_median"])
+   }
+ },
+   "dcast" = {
+     mean_poverty2 <- dcast(data, year ~ ., value.var = "income_median",
+                           fun.aggregate = mean)
+     names(mean_poverty2)[2] <- "poverty"
+   })

```

Unit: milliseconds

expr	min	lq	mean	median	uq	max	neval	cld
loop	1.340535	1.550222	2.539191	1.973241	2.796307	7.589922	100	a
dcast	1.646131	1.878245	2.897252	2.304728	3.075828	10.186388	100	a