

# Exercise: Functions

Day 4, Part B


```
> library(foreign)
> library(ggplot2)
> library(reshape2)
```

1. Write a function where the arguments are `in_file`, `x_variable`, `y_variable`, `facet_variable`, and `out_file`, that loads the specified input file, creates a scatter plot with the specified x, y, and faceting variables using `ggplot`, and then saves this plot to the specified output file while also returning the `ggplot` object. Assume the input file is a csv and the output file is a pdf. Hint: in this particular case, `aes_string()` is more user-friendly than `aes()` when calling `ggplot`.

```
> scatter_plot <- function(in_file, x_variable, y_variable, facet_variable, out_file) {
+   # load data
+   data <- read.csv(in_file)
+
+   # make plot
+   gg <- ggplot(data, aes_string(x = x_variable, y = y_variable)) +
+     facet_wrap(facet_variable) + geom_point()
+
+   # save plot
+   pdf(out_file)
+   print(gg)
+   dev.off()
+
+   # return plot
+   return(gg)
+ }
```

- a. Test this function using the data in `'data/gbd2015_global_deaths.csv'`, with year (`year_id`) on the x-axis, mean deaths (`val`) on the y-axis, and cause group (`cause_name`) as the faceting variable.


```
> main_dir <- "C:/Users/ngraetz/Documents/repos/r_training_penn/" # CHANGE TO YOUR LOCAL COPY
> scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
+   x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
+   out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'))
```



- b. Add assertions to the function to check that the input and output files are the expected format, and to provide a helpful error message if not. Test this by specifying an output file that is a jpeg, rather than a pdf. Hint: think back to the string functions lecture.

```
> scatter_plot <- function(in_file, x_variable, y_variable, facet_variable, out_file) {
+   # check input and output data format
+   if (!grepl('.csv$|.CSV$', in_file)) stop("'in_file' must be a csv file")
+   if (!grepl('.pdf$|.PDF$', out_file)) stop("'out_file' must be a pdf file")
+   # load data
+   data <- read.csv(in_file)
+   # make plot
+   gg <- ggplot(data, aes_string(x = x_variable, y = y_variable)) +
+     facet_wrap(facet_variable) + geom_point()
+   # save plot
+   pdf(out_file)
+   print(gg)
+   dev.off()
+   # return plot
+   return(gg)
+ }
>
> # this should work...
> scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
+   x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
```

```
+ out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'))
```



```
> # and this should throw an error...
> #scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
> #             x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
> #             out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.jpeg'))
```

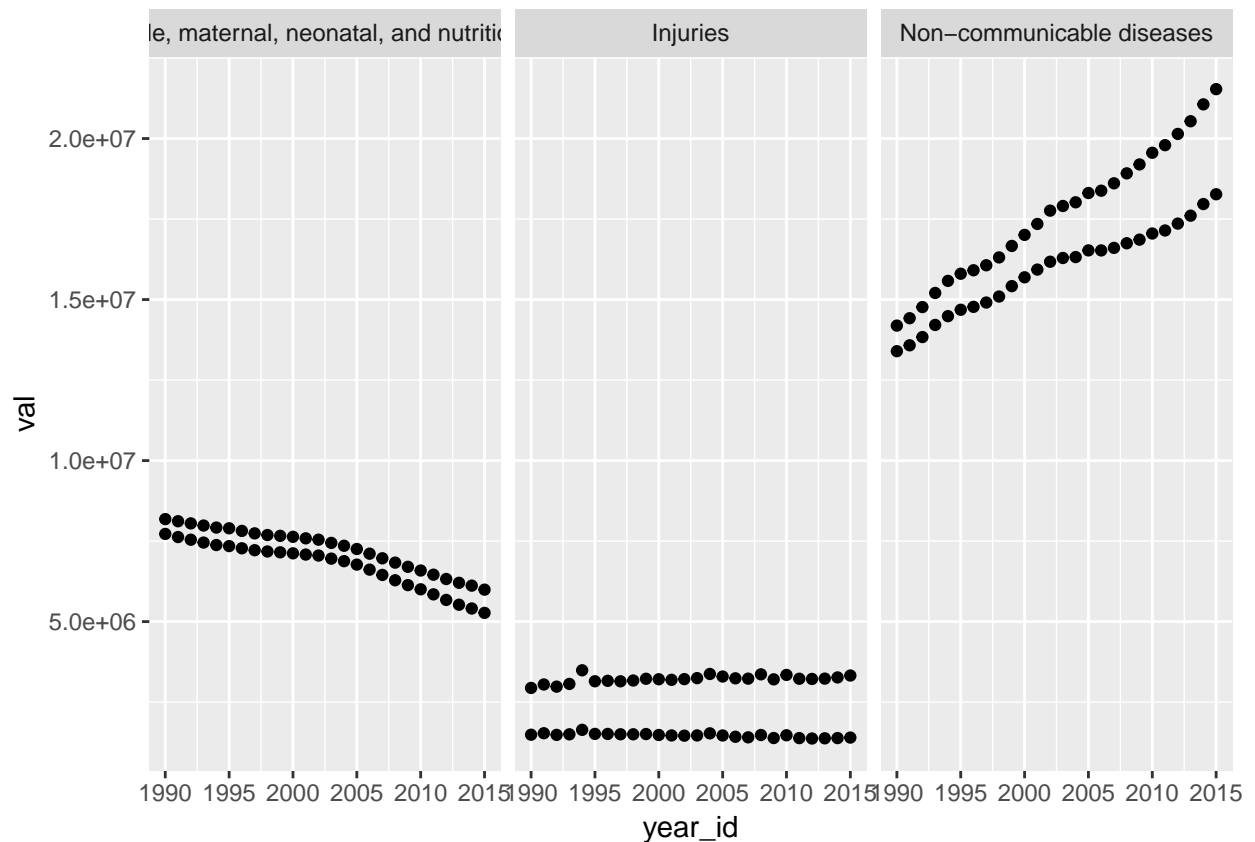
- c. Add assertions to the function to check that `x_variable`, `y_variable`, and `facet_variable` all exist as columns in the input data file, and to provide a helpful error message if not. Test this by providing an incorrect variable name.

```
> scatter_plot <- function(in_file, x_variable, y_variable, facet_variable, out_file) {
+   # check input and output data format
+   if (!grepl('.csv$|.CSV$', in_file)) stop("'in_file' must be a csv file")
+   if (!grepl('.pdf$|.PDF$', out_file)) stop("'out_file' must be a pdf file")
+   # load data
+   data <- read.csv(in_file)
+   # check that all variables exist
+   if (!x_variable %in% names(data)) stop(paste(x_variable, 'is not a column in', in_file))
+   if (!y_variable %in% names(data)) stop(paste(y_variable, 'is not a column in', in_file))
+   if (!facet_variable %in% names(data)) stop(paste(facet_variable, 'is not a column in', in_file))
+   # make plot
+   gg <- ggplot(data, aes_string(x = x_variable, y = y_variable)) +
+     facet_wrap(facet_variable) + geom_point()
+ }
```

```

+ # save plot
+ pdf(out_file)
+ print(gg)
+ dev.off()
+
+ # return plot
+ return(gg)
+ }
>
> # this should work...
> scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
+   x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
+   out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'))

```



```

> # and this should throw an error...
> #scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
> #   x_variable = 'year_id', y_variable = 'mean', facet_variable = 'cause_name',
> #   out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'))

```

- d. Add arguments `x_label`, `y_label`, and `title` to provide labels to the plot. Set them to default to “X Variable”, “Y Variable”, and “Title”, respectively. Test that this works as expected, both when values are supplied for these arguments and when they are not.

```


> scatter_plot <- function(in_file, x_variable, y_variable, facet_variable, out_file,
+   x_label = "X Variable", y_label = "Y Variable", title = "Title")
+
+ # check input and output data format
+ if (!grepl('.csv$|.CSV$', in_file)) stop("'in_file' must be a csv file")
+ if (!grepl('.pdf$|.PDF$', out_file)) stop("'out_file' must be a pdf file")

```

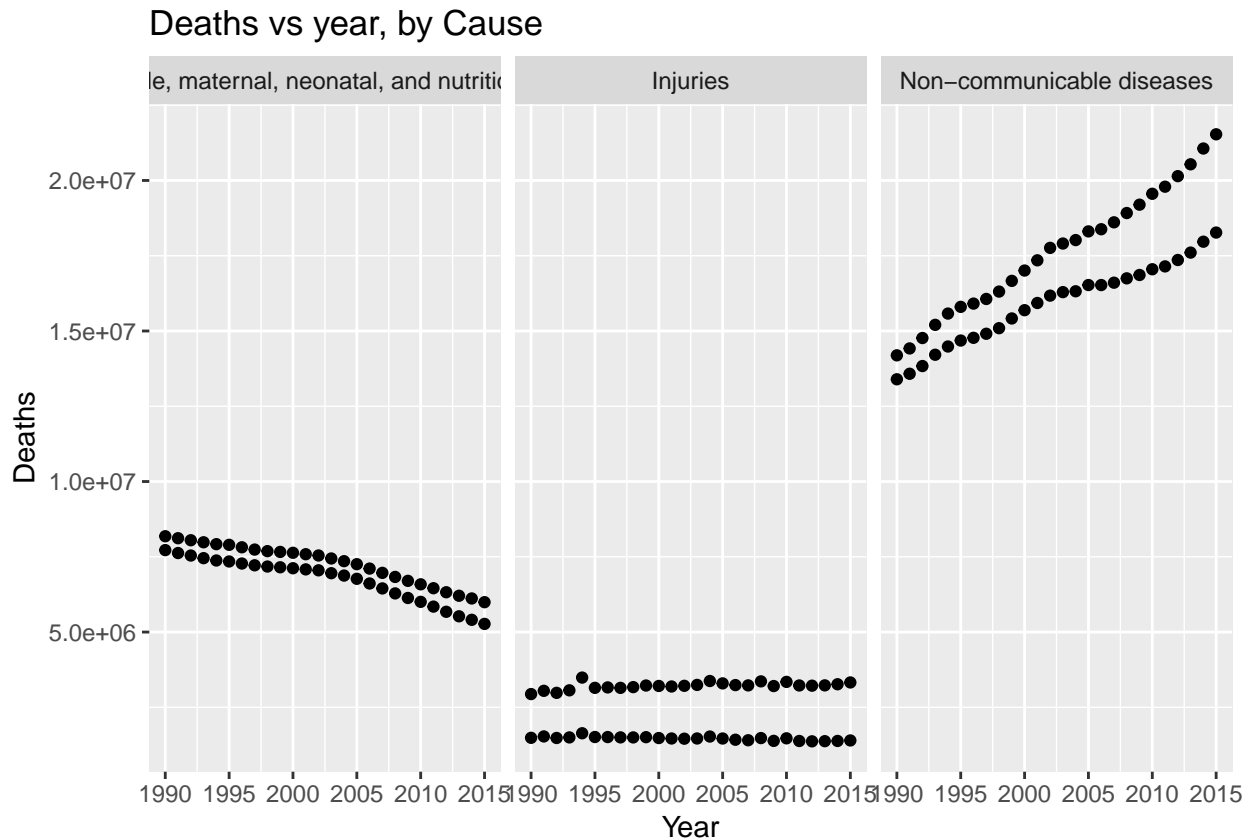
```

+
+ # load data
+ data <- read.csv(in_file)
+
+ # check that all variables exist
+ if (!x_variable %in% names(data)) stop(paste(x_variable, 'is not a column in', in_file))
+ if (!y_variable %in% names(data)) stop(paste(y_variable, 'is not a column in', in_file))
+ if (!facet_variable %in% names(data)) stop(paste(facet_variable, 'is not a column in', in_file))
+
+ # make plot
+ gg <- ggplot(data, aes_string(x = x_variable, y = y_variable)) +
+   facet_wrap(facet_variable) + geom_point() +
+   labs(x = x_label, y = y_label, title = title)
+
+ # save plot
+ pdf(out_file)
+ print(gg)
+ dev.off()
+
+ # return plot
+ return(gg)
+ }
+
+ > # using the defaults for labels/titles...
+ > scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
+   x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
+   out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'))

```



```
> # providing label/title arguments
> scatter_plot(in_file = paste0(main_dir, 'data/gbd2015_global_deaths.csv'),
+   x_variable = 'year_id', y_variable = 'val', facet_variable = 'cause_name',
+   out_file = paste0(main_dir, 'output/gbd2015_global_deaths_scatter.pdf'),
+   x_label = 'Year', y_label = 'Deaths', title = 'Deaths vs year, by Cause')
```



- Write a function that takes one argument `data`, a data frame, and returns a report of the number of missing and non-missing values for each variable, as well as the number of unique values, e.g.:

```
> # note that there are many ways to do this, and more than one valid output format
> missing_report <- function(data) {
+   report <- data.frame(vars = names(data),
+     missing = NA,
+     non_missing = NA,
+     values = NA)
+   for (v in names(data)) {
+     report[report$vars == v, 'missing'] <- sum(is.na(data[, v]))
+     report[report$vars == v, 'non_missing'] <- sum(!is.na(data[, v]))
+     report[report$vars == v, 'values'] <- length(unique(data[, v]))
+   }
+   return(report)
+ }
```

```
> data <- read.csv(paste0(main_dir, "data/ebola_polygon_data.csv"))
> missing_report(data)
vars missing non_missing values
1      UNIQ_ID      0      54      54
```

2	NAME	0	54	53
3	Country	0	54	6
4	Virus	0	54	4
5	CASE_TYPE	0	54	3
6	DATA_TYPE	0	54	1
7	LAT	0	54	46
8	LONG	0	54	45
9	SPR_ORDER	18	36	6
10	SOURCE_1	32	22	20
11	SOURCE_2	54	0	1
12	SOURCE_3	54	0	1
13	STR_DAY	42	12	11
14	STR_MNTH	18	36	10
15	STR_YEAR	4	50	15
16	END_DAY	52	2	3
17	END_MNTH	52	2	3
18	END_YEAR	52	2	3
19	REP_CASE	28	26	18
20	REP_DEATH	39	15	12
21	OB_ID	0	54	23
22	OB_STR_DAY	3	51	13
23	OB_STR_MNTH	0	54	10
24	OB_STR_YEAR	0	54	15
25	OB_END_DAY	13	41	12
26	OB_END_MNTH	0	54	11
27	OB_END_YEAR	0	54	14
28	OB_CASE	0	54	18
29	OB_DEATH	0	54	19

3. Repeat question 1 from lecture 4a using `lapply()` instead of a `for` loop.

```
> data <- lapply(1997:2015, function(year) {
+   if (year < 2004) {
+     sub <- read.csv(paste0(main_dir, "/data/wa_income_", year, ".csv"))
+   } else {
+     sub <- read.dta(paste0(main_dir, "/data/wa_income_", year, ".dta"))
+     sub <- plyr::rename(sub, c("FIPS" = "fips", "median_income" = "income_median"))
+   }
+   return(sub)
+ })
> data <- do.call('rbind', data)
> summary(data)
```

fips	year	income_median	poverty
Min. :53001	Min. :1997	Min. :27453	Min. : 6.60
1st Qu.:53019	1st Qu.:2001	1st Qu.:36992	1st Qu.:11.50
Median :53039	Median :2006	Median :42369	Median :14.10
Mean :53039	Mean :2006	Mean :43726	Mean :14.26
3rd Qu.:53059	3rd Qu.:2011	3rd Qu.:48693	3rd Qu.:16.40
Max. :53077	Max. :2015	Max. :81816	Max. :32.30

Bonus:

4. Repeat the first part of question 3 from lecture 4a using `sapply()` or `tapply()` instead of a `for` loop. Then do the same again, but calculate the ratio of the maximum to the minimum value instead of the mean.

```
> # mean with sapply...
> sapply(1997:2015, function(yy) {
+   mean(data[data$year == yy, 'income_median'])
+ })
[1] 36051.05 37450.33 37636.90 39212.51 38255.59 38989.44 39935.90 41323.79 41837.21
[10] 43795.28 46171.38 48219.13 47017.18 46621.33 47544.08 48141.85 49234.31 50873.38
[19] 52487.00
```

```
> # mean with tapply...
> tapply(data$income_median, data$year, mean)
1997      1998      1999      2000      2001      2002      2003      2004      2005
36051.05 37450.33 37636.90 39212.51 38255.59 38989.44 39935.90 41323.79 41837.21
      2006      2007      2008      2009      2010      2011      2012      2013      2014
43795.28 46171.38 48219.13 47017.18 46621.33 47544.08 48141.85 49234.31 50873.38
      2015
52487.00
```

```
> # max/min with sapply...
> sapply(1997:2015, function(yy) {
+   x <- data[data$year == yy, 'income_median']
+   max(x)/min(x)
+ })
[1] 1.868648 1.872010 1.816942 1.852431 1.857221 1.819798 1.761449 1.822441 1.842353
[10] 1.928604 1.906808 2.005580 1.945463 1.853792 1.985815 2.136076 2.010447 2.004605
[19] 2.034768
```

```
> # max/min with tapply...
> tapply(data$income_median, data$year, function(x) max(x) / min(x))
1997      1998      1999      2000      2001      2002      2003      2004      2005
1.868648 1.872010 1.816942 1.852431 1.857221 1.819798 1.761449 1.822441 1.842353
      2006      2007      2008      2009      2010      2011      2012      2013      2014
1.928604 1.906808 2.005580 1.945463 1.853792 1.985815 2.136076 2.010447 2.004605
      2015
2.034768
```

5. Load the US cigarette data ('data/us\_state\_cigarette\_data.rdata') and put all four of the data frames into a single list. Use the `Reduce()` function to merge all of these data sets in one go, rather than by calling `merge` directly three separate times. Do this once retaining only the rows where there are matches, and a second time retaining all rows.

```
> # load data and combine in a list
> load(paste0(main_dir, "/data/us_state_cigarette_data.rdata"), verbose = T)
Loading objects:
  cig_tax
  cig_csm
  pop
  locs

> cig_data_list <- list(cig_tax, cig_csm, pop, locs)
>
> # for only matching rows, we are totally fine with defaults
> cig_data <- Reduce(merge, cig_data_list)
> dim(cig_data)
[1] 204  6
```



```
> # for all rows, we need to set `all=T`, which requires defining a function that calls  
> # `merge()` with this setting  
> cig_data <- Reduce(function(x, y) merge(x, y, all = T), cig_data_list)  
> dim(cig_data)  
[1] 306 6
```